Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app

PROFILE

Forward-deployed AI engineer based in Honolulu — founder who builds AI agent platforms end to end. I embed with operators on ambiguous, high-stakes workflows and ship production agent systems: agent frameworks, orchestrators, MCP tool integrations, evaluation harnesses. Prototype in weeks, measure quality, productionize what carries load — with failure analysis and human-in-the-loop controls as the operating standard.

Three years running an evidence-grade agent control plane: typed state transitions, approval gates, audit receipts, recoverable runs, structured observability — a shared platform practice: one agent framework, reusable tools, and every deployment improving the next. AI-native: shipped AI-enabled products and workflows, with evals, release thresholds, and LLM-as-judge wired into delivery. Open to regular travel (~10–20%); not relocating off-island.

CORE COMPETENCIES

Engineering: Python · TypeScript · Node.js · SQL · Bash · REST APIs · Supabase · Postgres · Vercel · GitHub Actions (CI/CD)

AI & Agents: hands-on AI agent development & deployment (tool/API use, orchestration, state, failure handling) · agent frameworks on shared platforms · MCP protocol · LLM retrieval & structured outputs · LLM evaluation harnesses · prompt & context engineering

Forward Delivery: embedded discovery with non-technical operators · vague-problem-to-scoped-plan · pilot-in-weeks cadence · production rollout & adoption · AI platform feedback loops (field gaps → platform improvements) · human-in-the-loop controls

EXPERIENCE

Founder / Applied AI Systems Architect 2023 – Present

GlacierEQ — Honolulu, HI

- Founder-led platform practice: DISCOVER→FRAME→BUILD→INTEGRATE→EVALUATE→PROVE→DEPLOY with evidence receipts — embed on the operator's workflow, build the narrow end-to-end proof, productionize what carries load, and compile every deployment into reusable agents, skills, connectors, and playbooks.
- Hands-on agent framework delivery (anthropic-agent-coordinator): deterministic DAG scheduler, **95/95** tests green; tool calls proposed for review, never executed blind — human-in-the-loop control built into the orchestration layer; field learnings fed back into the shared framework.
- Shipped AI-enabled products and workflows measured by quality (job-app-helix): **989-test** agent control plane; evals with explicit release thresholds, gold sets and LLM-as-judge; evidence state machine (nominal / recoverable / hard-no-go) with the hard-failure path exiting non-zero and CI asserting it.
- Stateful retrieval and integration surface (unified-memory-mcp, mcp-iceberg-mansion): Mem0, Supermemory, Supabase pgvector, Pinecone, and Neo4j federated behind one MCP interface —

parallel fan-out, content-hash dedup; LLM retrieval-grounded generation with provenance and structured outputs.

- Reliable AI operations (apex-control-plane): cloud deployment, access controls, SHA-256 transport envelopes, idempotency keys, circuit breakers, dead-letter capture, fail-closed startup; failure analysis with root-cause quarantine instead of silent retries.
- Operate a **1,413-repository** applied-AI estate — **65 MCP connectors** integrated with claim tracking and idempotent automation; **20+ production repositories** maintained under test and CI; deployments measured on operator outcomes (time saved, errors cut, decisions improved).

Chief Inspector 2024

Diamond Head Home Inspections — Honolulu, HI

- Led residential inspection operations: field assessments, defect documentation, and client reporting — embedded delivery with non-technical customers; evidence-bound workflows where findings had to be right, documented, and defensible.

EDUCATION

University of Hawaii at Manoa 2016

Bachelor of Science

AWS Cloud Institute 2025

Cloud Computing program (brief attendance)

Colorado State University

Undergraduate coursework

SELECTED PROOF

Job-App Helix — attributable campaign engine

- Composes assessments into campaign-level **GO / NO-GO** with immutable evidence models, stable finding codes, and fail-closed decisions — evals with release thresholds and LLM-as-judge wired into CI.
- Recovery restricted to contingency evidence declared before reassessment — the engine never silently tunes values until they pass; audit receipts for every decision.

GlacierEQ estate — agent systems portfolio

- Durable patterns for claim tracking, evidence boundaries, audit receipts, observability with structured run states, and human-in-the-loop control across the full estate.
- Shared agent framework with a common data and context layer, reusable tools, and deployment tooling — field gaps returned to the platform so the next deployment needs less custom work.

SKILLS

Languages: Python · TypeScript · SQL · Bash · JavaScript / Node.js

AI / Agents: AI applications, agents, skills, connectors, workflow automations · multi-agent orchestration · RAG · LLM evaluation · prompt & context engineering

Data & Infra: Supabase · Postgres · Vercel · GitHub Actions · REST APIs · MCP tool integration patterns

Practice: evals & release thresholds · AI-native development (AI tools in the loop) · human-in-the-loop controls · observability · failure analysis · platform feedback loops · clear human-facing documentation

Evidence boundary: portfolio repositories are independent work. Domain-inspired demonstrations are not claims of company affiliation, production deployment at a customer's enterprise, grocery/retail depth, or certification. No Databricks/PySpark/dbt production depth claimed. No Afresh-platform production claims.