Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app

PROFILE

Coding agents are first-class users of the infrastructure I build. My agents read and run the surfaces I author: **65** MCP servers / connectors with typed contracts, docs, structured resource submission, evidence boundaries, and audit receipts. I write the prompts, skills, tool definitions, and agent instructions my own fleet runs on — and I measure what makes agents succeed or fail with evals, not vibes. Daily hands-on with agentic coding loops as the core build loop; clear technical writing that both humans and agents can act on.

Three years as a founder running an evidence-grade agent control plane from Honolulu: build the MCP server and agent integrations behind the product; make docs and code samples agent-ready; tune the prompts and tool definitions agents execute; measure success rate, time-to-first-working-call, and activation with real benchmarks. Speech-to-text pipeline experience through an estate-verified fal.ai + Whisper pipeline — integration-side, honestly bounded. Product sense for developer-facing work: I care whether the agent actually gets to a working integration.

CORE COMPETENCIES

Agent experience: agent-facing API & tool surface design · MCP servers & connectors · agent skills, prompts & instructions · tool definitions agents execute · machine-readable docs · llms.txt-style agent docs · quickstarts & SDK examples · developer-product sense

Eval & measurement: evaluation harnesses · agent benchmarks on real integration tasks · success-rate / time-to-first-call measurement · golden datasets · behavioral assertions · LLM-as-a-judge with written rubrics · trace intelligence · measure it, don't vibe it

AI & Voice: speech-to-text pipeline integration (Whisper-based, estate-verified) · LLM APIs · function calling · structured outputs · RAG · guardrails · fail-closed recovery

Production: JavaScript / TypeScript · Python (primary) · SQL · Bash · REST APIs · JSON Schema / Pydantic-style validation · GitHub Actions (CI/CD) · per-task latency & token-cost instrumentation · observability & run logging · clear technical writing

EXPERIENCE

Founder / Applied AI Systems Architect 2023 – Present

GlacierEQ — Honolulu, HI

- Build and maintain agent skills and integrations: author **65** MCP servers / connectors — the agent-readable layer with typed contracts, structured resource submission, and audit receipts — teaching agents how to integrate correctly the first time.
- Tune the prompts, tool definitions, and agent instructions my fleet executes: write the instructions agents run on, then measure them — behavioral assertions and golden datasets judging instruction quality, failure modes turned into better instructions.
- Measure it: eval harnesses that run real agents against real integration tasks; success-rate, time-to-first-working-call, and activation metrics with trace intelligence — evaluation before opinion.

- Make docs agent-ready: machine-readable documentation, quickstarts, and code samples written so an agent picks the right API, parameters, and patterns the first time — docs both audiences can act on.
- Use AI coding loops as a core part of how I work: daily hands-on with agentic coding as the build engine, with the curiosity about what makes agents succeed or fail that the Agent Experience discipline runs on.
- Run an evidence-grade agent control plane across a **1,413-repository** estate: guardrails, typed handoff contracts, fail-closed governance — backed by a **989-test** suite, **95/95** verification checks green.

EDUCATION

University of Hawaii at Manoa 2016
Bachelor of Science

AWS Cloud Institute 2025
Cloud Computing program

Colorado State University
Undergraduate coursework

SELECTED PROOF

MCP constellation — agent-facing infrastructure, 65 deep

- **65** authored MCP servers / connectors: governed API integrations with strict validation, structured resource submission, evidence boundaries, and audit receipts — the agent-facing surface this role builds.
- Direct evidence for the JD's nice-to-haves: working knowledge of MCP and LLM tool calling; shipped MCP servers that agents use in the wild.

Agent evals and speech pipeline, honestly bounded

- Eval harnesses measuring real agents on real integration tasks — the 'measure it' discipline the role centers on.
- fal.ai + Whisper pipeline (estate-verified) — speech-to-text pipeline experience, integration-side; no AssemblyAI-product, TTS, Realtime, or telephony-stack production claims.

SKILLS

Languages: JavaScript / TypeScript · Python · SQL · Bash · Node.js

Agent experience: MCP servers & connectors · agent skills & instructions · tool definitions · llms.txt-style machine-readable docs · quickstarts & code samples · developer-relations sensibility · technical writing

Eval & measurement: eval harnesses · agent benchmarks · golden datasets · behavioral assertions · LLM-as-a-judge · trace intelligence · success-rate & time-to-first-call measurement

AI & production: speech-to-text pipeline integration · LLM APIs · function calling · structured outputs · RAG · guardrails · CI regression gates · observability

Evidence boundary: portfolio repositories are independent work. No AssemblyAI-product, TTS/Realtime, or telephony-stack production claims. Whisper pipeline is integration-side, estate-verified. No named coding-agent endorsements — agentic coding loops, generically. GitHub is the forge of record.