Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app

PROFILE

Senior AI engineer building multi-agent frameworks for production: deterministic orchestration with explicit scheduling semantics, specialized-agent delegation, governed tool use, and evaluation infrastructure that gates what is allowed to run. Remote US — based in Honolulu.

Three years of intensive work on heavily constrained forensic-grade agent systems — adversarial conditions, evidence-grade accountability, zero margin for a wrong output. **95/95** on the Agent Coordinator's exact-revision test suite; **989-test** control plane with evals in dev and production; **65** governed MCP connectors, fail-closed.

CORE COMPETENCIES

Engineering: Python · TypeScript · SQL · Bash · REST APIs · Supabase · Postgres · GitHub Actions (CI/CD)

AI & Agents: multi-model LLM integration · multi-agent orchestration · agent workflow design · LLM evaluation frameworks · prompt & context engineering · RAG / search integration

Platform: privacy & secure data handling by construction · fail-closed tool authorization · audit receipts · idempotent automation

EXPERIENCE

Founder / Applied AI Systems Architect 2023 – Present

GlacierEQ — Honolulu, HI — Remote

- Design and operate a composable agent execution architecture: **1,413-repository** platform, agent governance, MCP integrations, context externalization, document workflows.
- Agent Coordinator — **95/95** recorded tests: deterministic multi-agent scheduling with ownership, dependencies, capacity, priority, shared budgets — human oversight as a state (proposed / executed / verified / blocked), not a policy doc.
- Evals as the floor: **989-test** control plane; behavior benchmarks in development and production, regression on every change — evals used to find and fix real failures before users do.
- Multi-model orchestration: integrate and route across models with capability, authentication, and authorization checks on **65** governed connectors — fail-closed by default, with audit receipts on every action.
- Ship MCP servers and agent skills connecting agents to real systems; partner-facing delivery — turn hard system behavior into decisions non-technical stakeholders use the same day.

EDUCATION

University of Hawaii at Manoa 2016

Bachelor of Science

AWS Cloud Institute 2025

Cloud Computing program

Colorado State University

Undergraduate coursework

SELECTED PROOF

Agent Coordinator — deterministic multi-agent scheduling

- **95/95** exact-revision tests: ownership, dependencies, capacity, priority, shared budgets, structured deferrals.
- Proposed / executed / verified / blocked states — the oversight semantics a multi-agent platform needs.

Eval & control plane — evaluation infrastructure

- **989-test** control plane with behavior benchmarks in dev and production and regression on every change.
- Evaluation harnesses that gate execution — evals used to find and fix real problems, not to decorate releases.

SKILLS

Languages: Python · TypeScript · SQL · Bash · JavaScript / Node.js

AI / Agents: multi-agent orchestration · LLM evaluation frameworks · multi-model integration · prompt & context engineering · RAG / search integration

Data & Infra: Supabase · Postgres · Vercel · GitHub Actions · REST APIs · MCP protocol

Practice: privacy & secure data handling · fail-closed authorization · audit receipts · verification gates

Evidence boundary: portfolio repositories are independent work. Orchestration and evaluation claims describe systems I designed and operate; they are not claims of employment, production deployment at another company, or certification.