Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app

PROFILE

I build agentic workflows the way this posting asks: ambiguous business problem in, well-framed workflow with clear success criteria and evaluation methodology out. Three years as a founder running an evidence-grade agent control plane from Honolulu — prototyping LLM-powered agents that reason, plan, and act across tools, APIs, and sensitive data sources, then hardening them to production with eval gates, typed handoff contracts, and fail-closed recovery. The bridge posture is native: I scope with operators, build the first-of-its-kind workflow, and leave audit receipts behind.

My integration layer is **65** authored MCP connectors — governed tool/API/data-source connections with typed contracts and structured validation on every boundary — plus evaluation harnesses with written rubrics (LLM-as-a-judge), golden datasets, and behavioral assertions. Agents must be reliable, observable, safe, and auditable from day one: that is my default operating posture, not a compliance add-on. Enterprise-customer depth is founder/operator (three years), honestly stated — no enterprise logos claimed.

CORE COMPETENCIES

Deployed delivery: embedded delivery posture · customer discovery · ambiguous-scope scoping · rapid prototyping → production · field friction → product requirements · technical-consulting posture · trusted builder relationships · handover discipline

Agentic AI: agentic workflows (reason / plan / act) · tool use & function calling · structured outputs · MCP servers & connectors · RAG pipelines · multi-agent loops · guardrails · fail-closed recovery · human-in-the-loop gates

Evaluation & quality: evaluation harnesses · golden datasets · behavioral assertions · LLM-as-a-judge with written rubrics · success-criteria framing · quality scoring · failure analysis · error taxonomies · trace intelligence · observability & run logging

Production: Python (primary) · TypeScript · SQL · Bash · REST APIs · JSON Schema / Pydantic-style validation · GitHub Actions (CI/CD) · Postgres · per-task latency & token-cost instrumentation · retry logic · documentation engineers actually read

EXPERIENCE

Founder / Applied AI Systems Architect 2023 – Present

GlacierEQ — Honolulu, HI

- Deliver the FDE motion end to end: translate ambiguous operator problems into well-framed agentic workflows with clear success criteria and written evaluation methodologies — prototype fast, measure honestly, harden what carries load.
- Ship LLM-powered agents that reason, plan, and act across tools, APIs, and data sources: typed handoff contracts govern every agent boundary; structured outputs and validation keep tool calls honest.
- Connect agents to the real world through **65** authored MCP connectors — governed integrations with typed contracts, evidence boundaries, and audit receipts on every tool boundary.

- Gate quality with evaluation harnesses before anything ships: golden datasets, behavioral assertions, LLM-as-a-judge with written rubrics — the eval methodology behind every production agent engagement.
- Design for failure in deployed settings: instrument the failure modes found in production — stuck loops, bad retrievals, tool misuse, context blowouts — and build the catchers: replay loops, fail-closed gates, recovery restricted to evidence declared before reassessment.
- Run an evidence-grade agent control plane across a **1,395-repository** estate: observability, trace intelligence, and recoverable runs — backed by a **989-test** suite, **62/62** verification checks green.

EDUCATION

University of Hawaii at Manoa 2016
Bachelor of Science

AWS Cloud Institute 2025
Cloud Computing program

Colorado State University
Undergraduate coursework

SELECTED PROOF

Job-App Helix — evaluation-gated campaign engine

- Attributable campaign engine with immutable evidence models and fail-closed **GO / NO-GO** decisions — the eval gate decides what carries load.
- Judge criteria with written rubrics behind every automated quality decision — the success-criteria discipline enterprise agent products depend on.

Governed integration layer — 65 authored MCP connectors

- Governed tool/API/data-source connections with typed contracts and structured validation on every boundary — the reusable integration layer behind every deployed agent engagement.
- Connectors span enterprise-style identity/access patterns: approval gates, access controls, audit receipts.

SKILLS

Languages: Python · TypeScript · SQL · Bash · JavaScript / Node.js

Agentic AI: agent workflows · tool/function calling · MCP · structured outputs · RAG · eval harnesses · LLM-as-a-judge · trace intelligence · guardrails · fail-closed recovery

Deployed engineering: discovery & scoping · rapid prototyping · embedded delivery · field-to-product feedback · handover discipline · technical writing

Production: REST APIs · Postgres · CI/CD (GitHub Actions) · JSON Schema validation · observability · latency & token-cost instrumentation · retry logic

Evidence boundary: portfolio repositories are independent work. No enterprise customer logos claimed; customer-facing delivery is founder/operator (three years), honestly stated. No Cohere-platform (North) internals claimed. No React-app-scale frontend production claims — TypeScript production is real; frontend flex is honest posture, not depth claimed. GitHub is the forge of record.