Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app

PROFILE

I build governed multi-agent systems that operate in regulated settings: a unified orchestration layer with typed handoff contracts, agent-assist tooling that keeps a human in the loop, and secure LLM-powered actions gated by evaluation before anything touches a real customer. Compliance isn't a review gate — it's the architecture.

Three years as a founder running an evidence-grade agent control plane from Honolulu: hand-rolled multi-agent loops with function calling, structured outputs, typed handoff contracts, durable recoverable runs, **65** MCP connectors, and full run logging so every action is reconstructable and auditable. Remote-first; experimentation-to-production fluency.

CORE COMPETENCIES

Engineering: TypeScript · Python · Node.js · SQL · Bash · REST APIs · containers · GitHub Actions (CI/CD) · Supabase · Postgres

AI & ML Systems: multi-agent orchestration · agent-assist systems · LLM-powered actions · eval harnesses & golden suites · trace intelligence · guardrails · fail-closed recovery · human-in-the-loop control · MCP protocol · prompt & context engineering

Production: regulated deployment · audit receipts & decision logs · CI regression gates · rollback as default posture · observability & run logging · experimentation discipline · documentation & handover

EXPERIENCE

Founder / Applied AI Systems Architect 2023 – Present

GlacierEQ — Honolulu, HI

- Build the unified orchestration layer for a multi-agent customer-support platform: hand-rolled multi-agent loops with function calling, structured outputs, and typed handoff contracts — every decision receipted, every action reconstructable from its trace.
- Gate LLM-powered actions by evaluation before they touch production: golden suites with behavioral assertions, judge criteria with written rubrics, CI regression gates that block a regression from merging. Fail-closed by default, never guess.
- Keep a human in the loop on consequential actions: approval gates typed into state transitions, recoverable runs with durable state that survives failure, rollback as the default posture — the governance pattern regulated deployments require.
- Run an evidence-grade agent control plane across a **1,413-repository** estate: human approval gates on consequential actions, typed state transitions, durable recoverable runs — backed by a **989-test** suite, **95/95** verification checks green.
- Integrate **65** MCP connectors — authoring MCP integrations that expose tools and structured resources to any MCP-compatible system — plus APIs, webhooks, and data pipelines, with evidence boundaries and audit receipts.

- Instrument quality, latency, and decision correctness across agent runs: per-task latency instrumentation and decision audit logs that make both benchmarking and post-incident forensics honest.

EDUCATION

University of Hawaii at Manoa 2016
Bachelor of Science

AWS Cloud Institute 2025
Cloud Computing program

Colorado State University
Undergraduate coursework

SELECTED PROOF

Job-App Helix — evaluation-gated campaign engine

- Attributable campaign engine with immutable evidence models, stable finding codes, and fail-closed **GO / NO-GO** decisions; recovery restricted to contingency evidence declared before reassessment.
- Verification gates on every agent action: no agent completion claim without independent provider readback — the eval gate decides, not the agent.

GlacierEQ estate — governed agent systems

- Durable patterns for claim tracking, evidence boundaries, audit receipts, and human-in-the-loop control across the full estate — the governance primitives a compliant-at-scale multi-agent platform needs.
- Per-task latency instrumentation and decision audit logs across agent runs — full reconstructability of every action.

SKILLS

Languages: TypeScript · Python · SQL · Bash · JavaScript / Node.js

AI / ML Systems: multi-agent orchestration · agent-assist systems · LLM-powered actions · eval harnesses & golden suites · trace intelligence · guardrails · fail-closed recovery · human-in-the-loop control · MCP protocol

Production: regulated deployment · audit receipts & decision logs · CI regression gates · rollback as default posture · observability & run logging · experimentation discipline · documentation & handover

Evidence boundary: portfolio repositories are independent work. Domain-inspired demonstrations are not claims of company affiliation, production deployment, or certification.