Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app

PROFILE

I build the platform harnesses production agentic workflows run on: orchestration layers with typed contracts, tool integrations that expose capabilities to any agent, shared context surfaces, and extension points other builders can trust and extend. My systems dogfood themselves — the agent platform I built runs my own operation daily.

Three years as a founder building governed agentic infrastructure from Honolulu: **65** MCP tool integrations, hand-rolled multi-agent loops with typed handoff contracts, eval gates that block regressions before production, and fail-closed guardrails — tracing, human-in-the-loop controls, safe rollout, rollback as default. TypeScript is my primary stack.

CORE COMPETENCIES

Engineering: TypeScript · Python · Node.js · REST APIs · webhooks & async fan-out · SQL · Bash · GitHub Actions (CI/CD) · containers · Supabase · Postgres

Agentic Platform: platform harnesses for agent builders · orchestration layers · tool-calling systems · MCP protocol · intent engineering · shared context surfaces · agent runtime patterns (job isolation, scheduling, secrets/access) · progressive delivery · safe rollout & rollback

Evals & Guardrails: eval harnesses & golden suites · regression detection gates · trace intelligence · human-in-the-loop control · fail-closed recovery · observability (metrics, logs, traces) · audit receipts & decision logs · telemetry-driven investment decisions

EXPERIENCE

Founder / Applied AI Systems Architect 2023 – Present

GlacierEQ — Honolulu, HI

- Design and operate platform harnesses for production agentic workflows: orchestration layers with typed handoff contracts, tool integrations, shared context surfaces, and extension points other teams can build on — every action receipted, every run reconstructable from its trace.
- Build **65** MCP tool integrations — authoring integrations that expose tools and structured resources to any MCP-compatible system, the tool-protocol layer agents and builders consume: evidence boundaries, audit receipts, and typed capability contracts on every connector.
- Gate agent behavior by evaluation: golden suites with behavioral assertions, judge criteria with written rubrics, CI regression gates that block a regression from merging — regression detection as infrastructure, not review. The forge itself fails closed on unapproved claims.
- Run guardrails for agentic systems others depend on: tracing across runs, human-in-the-loop approval gates typed into state transitions, safe rollout with rollback as the default posture, fail-closed recovery — a 1,398-repo estate backed by a **989-test** suite, **62/62** verification checks green.
- Ship TypeScript and Node.js production services with continuous-deployment discipline: automated pipelines, progressive delivery, monitoring and rollback, a bias toward frequent reversible changes — webhooks and async fan-out patterns that keep agent services loosely coupled.

- Dogfood the platform daily: the agent campaign engine I built runs my own operations end to end — success metrics tracked per run (adoption, quality, reliability), sharp edges surfaced from real usage turned into platform improvements.

EDUCATION

University of Hawaii at Manoa 2016
Bachelor of Science

AWS Cloud Institute 2025
Cloud Computing program

Colorado State University
Undergraduate coursework

SELECTED PROOF

Job-App Helix — evaluation-gated campaign engine

- Attributable campaign engine with immutable evidence models, stable finding codes, and fail-closed **GO / NO-GO** decisions; recovery restricted to contingency evidence declared before reassessment.
- Verification gates on every agent action: no agent completion claim without independent provider readback — the eval gate decides, not the agent.

GlacierEQ estate — governed agentic infrastructure

- Durable patterns for orchestration, tool integrations, claim tracking, evidence boundaries, audit receipts, and human-in-the-loop control across a 1,398-repo estate — the harness primitives a trusted production agent platform needs.
- Per-task latency instrumentation and decision audit logs across agent runs — full reconstructability of every action, the observability substrate safe rollout depends on.

SKILLS

Languages: TypeScript · Python · SQL · Bash · JavaScript / Node.js

Agentic Platform: platform harnesses · orchestration layers · tool-calling systems · MCP protocol · intent engineering · shared context · progressive delivery · safe rollout & rollback

Evals & Guardrails: eval harnesses & golden suites · regression detection · trace intelligence · human-in-the-loop control · fail-closed recovery · observability · audit receipts

Evidence boundary: portfolio repositories are independent work. Domain-inspired demonstrations are not claims of company affiliation, production deployment, or certification.