Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app

PROFILE

I build applied AI the way an evaluation engineer would: model evaluation harnesses, internal prototyping frameworks, and golden gates that decide whether a system ships. My last three years are spent on the machine that measures the AI — benchmarking, trace intelligence, and replayable eval loops.

Three years as a founder running an evidence-grade agent control plane from Honolulu: hand-rolled multi-agent loops, **65** MCP connectors, full run logging, a **989-test** suite with **62/62** verification checks green. TypeScript + Python production work across the estate. Remote-first; ships fast, measures everything.

CORE COMPETENCIES

Engineering: TypeScript · Python · Node.js · SQL · Bash · REST APIs · containers · GitHub Actions (CI/CD) · Supabase · Postgres

Applied AI: model evaluation · eval harnesses & golden suites · behavioral assertions · replay loops · trace intelligence · internal prototyping frameworks · prompt & context engineering · MCP protocol · function calling · structured outputs

Production: CI regression gates · fail-closed recovery · observability & run logging · latency & token-cost benchmarks · decision audit receipts · documentation & handover

EXPERIENCE

Founder / Applied AI Systems Architect 2023 – Present

GlacierEQ — Honolulu, HI

- Own model evaluation work end to end: golden suites with behavioral assertions, judge criteria with written rubrics, automated benchmarking loops — the harness that decides whether a model or agent change ships.
- Build internal prototyping frameworks: hand-rolled multi-agent loops with function calling, structured outputs, and typed handoff contracts — prototypes that run under the same eval discipline as production.
- Wire trace intelligence into every run: full logging of every prompt, tool call, and decision so any execution is reconstructable and debuggable from its trace; failed evals diagnosed from evidence, not argued about.
- Run an evidence-grade agent control plane across a **1,395-repository** estate: typed state transitions, human approval gates on consequential actions, durable state that survives a failed run, recoverable runs — backed by a **989-test** suite, **62/62** verification checks green.
- Integrate **65** MCP connectors — authoring MCP integrations that expose tools and structured resources to any MCP-compatible system — plus APIs, webhooks, and data pipelines, with evidence boundaries and audit receipts.
- TypeScript + Python production engineering across the estate: REST APIs, containers, Postgres/Supabase, CI/CD — the full stack behind the eval systems, not just the models.

EDUCATION

University of Hawaii at Manoa 2016
Bachelor of Science

AWS Cloud Institute 2025
Cloud Computing program

Colorado State University
Undergraduate coursework

SELECTED PROOF

Job-App Helix — evaluation-gated campaign engine

- Attributable campaign engine with immutable evidence models, stable finding codes, and fail-closed **GO / NO-GO** decisions; recovery restricted to contingency evidence declared before reassessment.
- Verification gates on every agent action: no agent completion claim without independent provider readback — the eval gate decides, not the agent.

GlacierEQ estate — applied AI systems

- Durable patterns for claim tracking, evidence boundaries, audit receipts, and human-in-the-loop control across the full estate.
- Per-task latency instrumentation and decision audit logs across agent runs — the tracing practice that makes benchmarking honest.

SKILLS

Languages: TypeScript · Python · SQL · Bash · JavaScript / Node.js

Applied AI: model evaluation · eval harnesses & golden suites · behavioral assertions · replay loops · trace intelligence · internal prototyping frameworks · prompt & context engineering · MCP protocol · function calling · structured outputs

Production: CI regression gates · fail-closed recovery · observability & run logging · latency & token-cost benchmarks · documentation & handover

Evidence boundary: portfolio repositories are independent work. Domain-inspired demonstrations are not claims of company affiliation, production deployment, or certification.