Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app

PROFILE

Senior AI platform engineer building the foundation agent teams build on: agent environments and actions, model selection and routing, context management, tool design, and evaluation infrastructure that drives task-completion quality — evals as the iteration loop, not the afterthought.

Three years of intensive work on a heavily constrained forensic-grade agent system — adversarial conditions, evidence-grade accountability, zero margin for a wrong output. **95/95** exact-revision tests on the Agent Coordinator; **989-test** control plane; **65** governed MCP connectors, fail-closed.

The domain isn't abstract: three years inside a live family-court matter as the pro se party — long records, zero margin for a wrong citation. I know what legal AI gets wrong from the user's side, and for me reliability is not abstract.

CORE COMPETENCIES

Engineering: Python · TypeScript · SQL · Bash · REST APIs · Supabase · Postgres · GitHub Actions (CI/CD)

AI & Agents: agent environments & actions · model selection & routing · context-window management · tool design · practical evaluation harnesses · multi-agent orchestration · RAG / multi-source retrieval

Platform: evaluation infrastructure that gates execution · fail-closed authorization · audit receipts · idempotent automation

EXPERIENCE

Founder / Applied AI Systems Architect 2023 – Present

GlacierEQ — Honolulu, HI

- Build the systems that make AI agents dependable for high-stakes professional work: agent environments, actions, and tools designed against real workflows — not demos.
- Practical evaluations that capture what 'excellent' means: **989-test** control plane with behavior benchmarks in dev and production; evals drive the iteration loop and unlock new capabilities.
- Agent Coordinator — **95/95** recorded tests: model-agnostic deterministic scheduling with ownership, dependencies, capacity, priority, shared budgets; human oversight by construction (proposed / executed / verified / blocked).
- Context and session state done right: durable agent state with inspectable transitions, bounded actions, recoverable workflows — state that survives failure and shows its work.
- **65** governed MCP connectors with capability, authentication, and authorization checks, fail-closed by default — the tool layer agent workflows depend on.

EDUCATION

University of Hawaii at Manoa 2016

Bachelor of Science

AWS Cloud Institute 2025

Cloud Computing program

Colorado State University
Undergraduate coursework

SELECTED PROOF

Eval-driven agent iteration

- Evaluation harnesses used to find and fix real problems — retrieval quality, task-completion rates, latency — before users do.
- Representative pattern: multi-source retrieval agents tuned against practical evals until task completion, not just model scores, moves.

Agent infrastructure that gates what runs

- **989-test** control plane; regression on every change; approval gates on consequential actions.
- Fail-closed tool authorization with audit receipts — trust as a business enabler, built into the platform.

SKILLS

Languages: Python · TypeScript · SQL · Bash · JavaScript / Node.js

AI / Agents: agent environments & actions · model selection & routing · context management · tool design · evaluation harnesses · multi-agent orchestration · RAG / retrieval

Data & Infra: Supabase · Postgres · Vercel · GitHub Actions · REST APIs · MCP protocol

Practice: evidence-grade accountability · fail-closed authorization · audit receipts · plain-language delivery to non-technical stakeholders

Evidence boundary: portfolio repositories are independent work. Domain-inspired demonstrations are not claims of company affiliation, production deployment, or certification. Legal-systems perspective comes from firsthand pro se experience, not legal employment.