Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app

PROFILE

Agent-systems architect based in Honolulu — I govern how multi-agent systems are allowed to ship: deterministic orchestration, fail-closed tool handoffs, and evaluation frameworks that decide GO / NO-GO before production. Three years running evidence-grade agent infrastructure: typed state transitions, approval gates, recoverable runs.

My production agent code is hand-rolled — I work the patterns the frameworks name (LangGraph checkpoint concepts are adopted as design discipline in my control-plane docs; portfolio repos ai-pdf-chatbot-langchain and apex-memory-agent touch LangChain/LangGraph directly). Embedded delivery: scope with the people whose workflow it is, build the narrow end-to-end proof, harden what carries load. Open to 20% travel; not relocating off-island.

CORE COMPETENCIES

Engineering: Python · TypeScript · Node.js · SQL · Bash · REST APIs · Supabase · Postgres · Vercel · GitHub Actions (CI/CD)

AI & Agents: multi-agent orchestration (deterministic, hand-rolled) · agent evaluation frameworks · rubric-based scoring · prompt engineering & optimization · short-term / long-term memory patterns · vector stores · RAG · MCP tool integration · error handling

Deployment: Linux · containers · CI/CD pipelines · trace-logging observability · idempotent automation · networking & security fundamentals

Practice: customer-embedded scoping · technical assessments · consultative delivery · audit receipts · human-facing technical documentation

EXPERIENCE

Founder / Applied AI Systems Architect 2023 – Present

GlacierEQ — Honolulu, HI

- Design multi-agent system architectures with deterministic coordination: the agent coordinator enforces assignment-bound, fail-closed tool-proposal handoffs bound to plan SHA-256 — **95/95** behavioral checks green (hash-bound candidate receipt).
- Design and operate comprehensive evaluation frameworks: the campaign engine composes assessments into **GO / NO-GO** with immutable evidence models and stable finding codes; recovery is restricted to contingency evidence declared before reassessment — the engine never silently tunes values until they pass.
- Ship rubric-based scoring in production code: honesty/novelty scoring that explicitly excludes self-scored boilerplate — evaluation of the evaluator, not just of the agent.
- Integrated **65 MCP connectors** as the governed tool-integration layer — tool integration, API design, and error-handling patterns at the boundary between customer environments and the platform.
- Operate two agent platforms under heavy test: a **989-test** campaign engine and a **988-test** control plane — typed state transitions, approval gates, recoverable runs, full trace logging, regression discipline.

- Portfolio repos reach the LangChain/LangGraph ecosystem directly (ai-pdf-chatbot-langchain, pushed 2026-09-13; LangGraph checkpoint concepts adopted in the control plane's upstream adoption matrix) — framework-adjacent patterns in hand-rolled production code.
- Customer-embedded delivery practice: map the operator's workflow, assess the technical reality, present recommendations non-technical stakeholders act on the same day.

EDUCATION

University of Hawaii at Manoa 2016
Bachelor of Science

AWS Cloud Institute 2025
Cloud Computing program

Colorado State University
Undergraduate coursework

SELECTED PROOF

Job-App Helix — attributable campaign engine

- Composes assessments into campaign-level **GO / NO-GO** with immutable evidence models, stable finding codes, and fail-closed decisions.
- Recovery restricted to contingency evidence declared before reassessment — the engine never silently tunes values until they pass.

Anthropic Agent Coordinator — deterministic multi-agent orchestration

- Assignment-bound tool proposals: only fully assigned tasks may emit a proposal; deferred, unknown, duplicate, oversized, or malformed proposals fail closed.
- Hand-rolled orchestration — the LangGraph-named patterns (checkpoints, state transitions) implemented as native code with typed state.

Genius engine — rubric-based scoring

- Novelty/honesty scoring judges problem/cause/mechanism only, explicitly excluding implementation-boilerplate self-scoring.

SKILLS

Languages: Python · TypeScript · SQL · Bash · JavaScript / Node.js

AI / Agents: multi-agent orchestration · agent evaluation frameworks · rubric scoring · prompt engineering & optimization · memory patterns · RAG · MCP protocol

Data & Infra: Supabase · Postgres · Vercel · GitHub Actions · REST APIs · Linux · containers · observability discipline

Practice: verification gates · audit receipts · idempotent automation · consultative delivery · technical assessments

Evidence boundary: portfolio repositories are independent work. Domain-inspired demonstrations are not claims of company affiliation, production deployment at customer scale, or certification.