Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app

PROFILE

Forward-deployed AI engineer and systems architect based in Honolulu. I set the technical direction for applied AI systems end to end — agents, skills, connectors, workflow automations, evaluation harnesses — across a 1,413-repository applied-AI estate, and ship production workflows with evals, review gates, human-in-the-loop controls, and governance as the operating standard.

Leadership through architecture and practice: I define reusable platform patterns (claim tracking, evidence boundaries, audit receipts, structured run-state observability) that every project in the estate inherits; mentor through design and code review across technical surfaces; and return platform requirements to shared infrastructure — the compounding loop Reddit's AI Enablement practice describes.

Three years intensive applied-AI systems work (2023–present). Embedding with operators on ambiguous, high-stakes workflows; prototyping in days, productionizing within weeks; transferring maintainable workflows to owners. Open to travel; not relocating off-island.

CORE COMPETENCIES

Engineering: Python · TypeScript · Node.js · SQL · Bash · REST APIs · Supabase · Postgres · Vercel · GitHub Actions (CI/CD)

AI & Agents: hands-on AI agent development & deployment (tool/API use, orchestration, state, failure handling) · MCP protocol · reusable tool integrations · LLM retrieval & structured outputs · LLM evaluation harnesses · prompt & context engineering

Technical Leadership: architecture direction across multi-system estates · design & code review mentoring · reusable platform patterns adopted estate-wide · delivery frameworks (DISCOVER→FRAME→BUILD→INTEGRATE→EVALUATE→PROVE→DEPLOY) · investment priorities tied to measurable quality/speed/capacity/cost gains

Internal Delivery: embedded scoping with non-technical stakeholders · evals & review gates · human-in-the-loop controls · observability & failure analysis · trusted production workflows with governance, support, reversibility, and named owners

EXPERIENCE

Founder / Applied AI Systems Architect 2023 – Present

GlacierEQ — Honolulu, HI

- Set the vision and technical direction for applied AI across a **1,413-repository** estate: governed agent systems with policy enforcement, approval gates, audit receipts, and structured run-state observability — platform patterns every project inherits.
- Forward-deployed delivery framework as executable code (Forward-Deployed-Agentic-Powerhouse): DISCOVER→FRAME→BUILD→INTEGRATE→EVALUATE→PROVE→DEPLOY with evidence receipts — a repeatable practice for turning ambiguous operator problems into production workflows, then transferring maintainable ownership.

- Hands-on agent development and deployment (anthropic-agent-coordinator): deterministic DAG scheduler, **95/95** tests green; tool calls proposed for review, never executed blind — human-in-the-loop control built into the orchestration layer.
- Evaluation-first delivery (job-app-helix): **989-test** agent control plane; evals with explicit release thresholds, gold sets and LLM-as-judge; evidence state machine (nominal / recoverable / hard-no-go) with the hard-failure path exiting non-zero and CI asserting it.
- Reliable AI operations (apex-control-plane): cloud deployment, access controls, SHA-256 transport envelopes, idempotency keys, circuit breakers, dead-letter capture, fail-closed startup; failure analysis with root-cause quarantine instead of silent retries.
- Stateful retrieval (unified-memory-mcp, mcp-iceberg-mansion): Mem0, Supermemory, Supabase pgvector, Pinecone, and Neo4j federated behind one MCP interface — parallel fan-out, content-hash dedup; LLM retrieval-grounded generation with provenance and structured outputs.
- **65 MCP connectors** integrated with claim tracking and idempotent automation; reusable agents, skills, connectors, and playbooks compiled as the compounding output; **20+ production repositories** maintained under test and CI.

Chief Inspector 2024

Diamond Head Home Inspections — Honolulu, HI

- Led residential inspection operations: field assessments, defect documentation, and client reporting — embedded delivery with non-technical stakeholders; business judgment from regulated, evidence-bound work.

EDUCATION

University of Hawaii at Manoa 2016

Bachelor of Science

AWS Cloud Institute 2025

Cloud Computing program (brief attendance)

Colorado State University

Undergraduate coursework

SELECTED PROOF

Job-App Helix — attributable campaign engine

- Composes assessments into campaign-level **GO / NO-GO** with immutable evidence models, stable finding codes, and fail-closed decisions — review gates with real veto power.
- Recovery restricted to contingency evidence declared before reassessment — the engine never silently tunes values until they pass; reversibility and named accountability by design.

GlacierEQ estate — agent systems portfolio

- Durable patterns for claim tracking, evidence boundaries, audit receipts, observability with structured run states, and human-in-the-loop control — adopted across the full estate, not one project.
- Connects multiple technical surfaces — agents, APIs, databases, cloud — into one coherent, verifiable working system; 65 MCP connectors as the integration surface.

SKILLS

Languages: Python · TypeScript · SQL · Bash · JavaScript / Node.js

AI / Agents: AI applications, agents, skills, connectors, workflow automations · multi-agent orchestration · RAG · LLM evaluation · prompt & context engineering

Data & Infra: Supabase · Postgres · Vercel · GitHub Actions · REST APIs · MCP tool integration patterns · enterprise identity and access patterns

Practice: evals · review gates · human-in-the-loop controls · observability · failure analysis · governance & reversibility · reusable platform patterns · clear human-facing documentation

Evidence boundary: portfolio repositories are independent work. Domain-inspired demonstrations are not claims of company affiliation, production deployment at a customer's enterprise, or certification. Leadership claims describe architecture direction, design/code review mentoring, and practice-setting across his own estate — no formal people-management or corporate reporting lines claimed. No Reddit-platform production depth claimed.