# Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app

## PROFILE

Forward-deployed AI engineer based in Honolulu. I embed with operators on ambiguous, high-stakes workflows and ship production AI systems end to end — agents, skills, connectors, workflow automations, evaluation harnesses — prototyping in days and productionizing within weeks, with evals, review gates, and human-in-the-loop controls as the operating standard.

Three years running an evidence-grade agent control plane: typed state transitions, approval gates, audit receipts, recoverable runs, structured observability. Every workflow ships with named ownership, reversibility, and a transferable operating handoff — internal stakeholders own it after delivery. Open to travel; not relocating off-island.

## CORE COMPETENCIES

**Engineering:** Python · TypeScript · Node.js · SQL · Bash · REST APIs · Supabase · Postgres · Vercel · GitHub Actions (CI/CD)

**AI & Agents:** hands-on AI agent development & deployment (tool/API use, orchestration, state, failure handling) · MCP protocol · LLM retrieval & structured outputs · LLM evaluation harnesses · prompt & context engineering

**Internal Delivery:** embedded scoping with non-technical stakeholders · prototype-in-days cadence · evals & review gates · human-in-the-loop controls · observability & failure analysis · trusted production workflows with governance, support, reversibility, and named owners

## EXPERIENCE

**Founder / Applied AI Systems Architect** 2023 – Present

*GlacierEQ — Honolulu, HI*

- Forward-deployed motion as executable code (Forward-Deployed-Agentic-Powerhouse): DISCOVER→FRAME→BUILD→INTEGRATE→EVALUATE→PROVE→DEPLOY with evidence receipts — map the operator's workflow, build the narrow end-to-end proof, productionize what carries load, transfer maintainable ownership.
- Stateful retrieval (unified-memory-mcp, mcp-iceberg-mansion): Mem0, Supermemory, Supabase pgvector, Pinecone, and Neo4j federated behind one MCP interface — parallel fan-out, content-hash dedup; LLM retrieval-grounded generation with provenance and structured outputs.
- Hands-on agent development and deployment (anthropic-agent-coordinator): deterministic DAG scheduler, **95/95** tests green; tool calls proposed for review, never executed blind — human-in-the-loop control built into the orchestration layer.
- Evaluation-first delivery (job-app-helix): **989-test** agent control plane; evals with explicit release thresholds, gold sets and LLM-as-judge; evidence state machine (nominal / recoverable / hard-no-go) with the hard-failure path exiting non-zero and CI asserting it.
- Reliable AI operations (apex-control-plane): cloud deployment, access controls, SHA-256 transport envelopes, idempotency keys, circuit breakers, dead-letter capture, fail-closed startup; failure analysis with root-cause quarantine instead of silent retries.

- Operate a **1,413-repository** applied-AI estate — **65 MCP connectors** integrated with claim tracking and idempotent automation; reusable agents, skills, connectors, and playbooks compiled as the compounding output; **20+ production repositories** maintained under test and CI.

**Chief Inspector** 2024

*Diamond Head Home Inspections — Honolulu, HI*

- Led residential inspection operations: field assessments, defect documentation, and client reporting — embedded delivery with non-technical stakeholders, and measurable workflow judgment from regulated, evidence-bound work.

## EDUCATION

**University of Hawaii at Manoa** 2016

*Bachelor of Science*

**AWS Cloud Institute** 2025

*Cloud Computing program (brief attendance)*

**Colorado State University**

*Undergraduate coursework*

## SELECTED PROOF

**Job-App Helix — attributable campaign engine**

- Composes assessments into campaign-level **GO / NO-GO** with immutable evidence models, stable finding codes, and fail-closed decisions — review gates with real veto power.
- Recovery restricted to contingency evidence declared before reassessment — the engine never silently tunes values until they pass; reversibility and named accountability by design.

**GlacierEQ estate — agent systems portfolio**

- Durable patterns for claim tracking, evidence boundaries, audit receipts, observability with structured run states, and human-in-the-loop control across the full estate.
- Connects multiple technical surfaces — agents, APIs, databases, cloud — into one coherent, verifiable working system; 65 MCP connectors as the integration surface.

## SKILLS

**Languages:** Python · TypeScript · SQL · Bash · JavaScript / Node.js

**AI / Agents:** AI applications, agents, skills, connectors, workflow automations · multi-agent orchestration · RAG · LLM evaluation · prompt & context engineering

**Data & Infra:** Supabase · Postgres · Vercel · GitHub Actions · REST APIs · MCP tool integration patterns

**Practice:** evals · review gates · human-in-the-loop controls · observability · failure analysis · governance & reversibility · clear human-facing documentation

*Evidence boundary: portfolio repositories are independent work. Domain-inspired demonstrations are not claims of company affiliation, production deployment at a customer's enterprise, or certification. No Reddit-platform production depth claimed.*