# Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app

## PROFILE

I build the production machinery that lets entity agents ship: automated benchmarking and evaluation, golden gates, evidence adapters, replayable run stores, and instrumentation across cost, quality, latency, and health — the deployment discipline that decides whether an agent is allowed to run. Rollback isn't a feature; it's the default. Fail-closed by default, never guess.

Three years as a founder running an evidence-grade agent control plane from Honolulu: hand-rolled multi-agent loops with function calling, structured outputs, typed handoff contracts, durable recoverable runs, **65** MCP connectors, and full run logging so any execution can be reconstructed from its trace. I own eval, deployment, and recovery end to end — principal scope carried as founder experience, not invented title.

## CORE COMPETENCIES

**Engineering:** TypeScript · Python · Node.js · SQL · Bash · REST APIs · containers · GitHub Actions (CI/CD) · Supabase · Postgres

**AI & Evaluation:** automated benchmarking · evaluation harnesses · golden gates · behavioral assertions · evidence adapters · replayable run stores · trace intelligence · cost/quality/latency/health instrumentation · hand-rolled multi-agent loops · guardrails · fail-closed recovery · MCP protocol

**Production:** agent-platform deployment · rollback mechanisms · CI regression gates · decision audit receipts · observability & run logging · experimentation discipline · documentation & handover

## EXPERIENCE

**Founder / Applied AI Systems Architect** 2023 – Present

*GlacierEQ — Honolulu, HI*

- Ship the production machinery for AI-native agents: automated benchmarking and evaluation suites that gate what deploys — golden suites with behavioral assertions, judge criteria with written rubrics — so no agent change ships without measured evidence it holds.
- Run evidence adapters and replayable run stores: full logging of every prompt, tool call, and decision so any run can be reconstructed from its trace, failures diagnosed from evidence, and regressions replayed on demand.
- Instrument cost, quality, latency, and health across agent runs: per-task latency instrumentation and decision audit logs that make benchmarking honest and rollback decisions data-driven.
- Hand-roll multi-agent loops with function calling, structured outputs, and typed handoff contracts: orchestration the operator can audit, every decision receipted — durable state that survives a failed run, recoverable runs.
- Run an evidence-grade agent control plane across a **1,413-repository** estate: typed state transitions, human approval gates on consequential actions — backed by a **989-test** suite, **95/95** verification checks green.

- Integrate **65** MCP connectors — authoring MCP integrations that expose tools and structured resources to any MCP-compatible system — plus APIs, webhooks, and data pipelines, with evidence boundaries and audit receipts.

## EDUCATION

**University of Hawaii at Manoa** 2016
*Bachelor of Science*

**AWS Cloud Institute** 2025
*Cloud Computing program*

**Colorado State University**
*Undergraduate coursework*

## SELECTED PROOF

**Job-App Helix — evaluation-gated campaign engine**

- Attributable campaign engine with immutable evidence models, stable finding codes, and fail-closed **GO / NO-GO** decisions; recovery restricted to contingency evidence declared before reassessment.
- Verification gates on every agent action: no agent completion claim without independent provider readback — the eval gate decides, not the agent.

**GlacierEQ estate — production agent systems**

- Durable patterns for claim tracking, evidence boundaries, audit receipts, and human-in-the-loop control across the full estate.
- Per-task latency instrumentation and decision audit logs across agent runs — the tracing practice that makes benchmarking honest.

## SKILLS

**Languages:** TypeScript · Python · SQL · Bash · JavaScript / Node.js

**AI / Evaluation:** automated benchmarking · evaluation harnesses · golden gates · replay loops · trace intelligence · cost/quality/latency instrumentation · guardrails · MCP protocol · typed handoff contracts · prompt & context engineering

**Production:** agent-platform deployment · rollback mechanisms · CI regression gates · fail-closed recovery · observability & run logging · experimentation discipline · documentation & handover

*Evidence boundary: portfolio repositories are independent work. Domain-inspired demonstrations are not claims of company affiliation, production deployment, or certification.*