Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app

PROFILE

I build the production machinery that lets entity agents ship: automated benchmarking and evaluation, golden gates, evidence adapters, replayable run stores, and instrumentation across cost, quality, latency, and health — the deployment discipline that decides whether an agent is allowed to run. Rollback isn't a feature; it's the default. Fail-closed by default, never guess.

Three years as a founder running an evidence-grade agent control plane from Honolulu: hand-rolled multi-agent loops with function calling, structured outputs, typed handoff contracts, durable recoverable runs, **65** MCP connectors, and full run logging so any execution can be reconstructed from its trace. I own eval, deployment, and recovery end to end — principal scope carried as founder experience, not invented title.

CORE COMPETENCIES

Engineering: TypeScript · Python · Node.js · SQL · Bash · REST APIs · containers · GitHub Actions (CI/CD) · Supabase · Postgres

AI & Evaluation: automated benchmarking · evaluation harnesses · golden gates · behavioral assertions · evidence adapters · replayable run stores · trace intelligence · cost/quality/latency/health instrumentation · hand-rolled multi-agent loops · guardrails · fail-closed recovery · MCP protocol

Production: agent-platform deployment · rollback mechanisms · CI regression gates · decision audit receipts · observability & run logging · experimentation discipline · documentation & handover

EXPERIENCE

Founder / Applied AI Systems Architect 2023 – Present

GlacierEQ — Honolulu, HI

- Ship the production machinery for AI-native agents: automated benchmarking and evaluation suites that gate what deploys — golden suites with behavioral assertions, judge criteria with written rubrics — so no agent change ships without measured evidence it holds.
- Run evidence adapters and replayable run stores: full logging of every prompt, tool call, and decision so any run can be reconstructed from its trace, failures diagnosed from evidence, and regressions replayed on demand.
- Instrument cost, quality, latency, and health across agent runs: per-task latency instrumentation and decision audit logs that make benchmarking honest and rollback decisions data-driven.
- Hand-roll multi-agent loops with function calling, structured outputs, and typed handoff contracts: orchestration the operator can audit, every decision receipted — durable state that survives a failed run, recoverable runs.
- Run an evidence-grade agent control plane across a **1,395-repository** estate: typed state transitions, human approval gates on consequential actions — backed by a **989-test** suite, **62/62** verification checks green.

- Integrate **65** MCP connectors — authoring MCP integrations that expose tools and structured resources to any MCP-compatible system — plus APIs, webhooks, and data pipelines, with evidence boundaries and audit receipts.

EDUCATION

University of Hawaii at Manoa 2016
Bachelor of Science

AWS Cloud Institute 2025
Cloud Computing program

Colorado State University
Undergraduate coursework

SELECTED PROOF

Job-App Helix — evaluation-gated campaign engine

- Attributable campaign engine with immutable evidence models, stable finding codes, and fail-closed **GO / NO-GO** decisions; recovery restricted to contingency evidence declared before reassessment.
- Verification gates on every agent action: no agent completion claim without independent provider readback — the eval gate decides, not the agent.

GlacierEQ estate — production agent systems

- Durable patterns for claim tracking, evidence boundaries, audit receipts, and human-in-the-loop control across the full estate.
- Per-task latency instrumentation and decision audit logs across agent runs — the tracing practice that makes benchmarking honest.

SKILLS

Languages: TypeScript · Python · SQL · Bash · JavaScript / Node.js

AI / Evaluation: automated benchmarking · evaluation harnesses · golden gates · replay loops · trace intelligence · cost/quality/latency instrumentation · guardrails · MCP protocol · typed handoff contracts · prompt & context engineering

Production: agent-platform deployment · rollback mechanisms · CI regression gates · fail-closed recovery · observability & run logging · experimentation discipline · documentation & handover

Evidence boundary: portfolio repositories are independent work. Domain-inspired demonstrations are not claims of company affiliation, production deployment, or certification.