Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app

PROFILE

I build the trust machinery behind every system I ship: eval harnesses with replay loops that validate prompt and configuration changes before anything touches production, guardrails with typed state transitions and approval gates on consequential actions, per-run monitoring — and a defined 'what happens when it's unsure' path: the agent halts and escalates, never guesses. Fail-closed by default.

Three years operating an evidence-grade agent control plane from Honolulu: hand-rolled multi-agent loops with function calling, structured outputs, typed handoff contracts, retrieval over my own estate, and 65 MCP connectors — delivered with audit receipts and evidence boundaries. Remote-first; client-embedded.

CORE COMPETENCIES

Engineering: Python · TypeScript · Node.js · SQL · Bash · REST APIs · webhooks · Supabase · Postgres · Vercel · GitHub Actions (CI/CD) · containers

AI & Agents: LLM evaluation harnesses · replay loops · guardrails · fail-closed recovery · LLM agent frameworks (hand-rolled multi-agent loops) · function calling · structured outputs · RAG · typed handoff contracts · prompt & context engineering · MCP protocol · human-in-the-loop control

Deployment: client-embedded scoping · demos & POCs · monitoring & observability · latency & token-cost benchmarks · verification gates · audit receipts · idempotent automation · evidence boundaries

EXPERIENCE

Founder / Applied AI Systems Architect 2023 – Present

GlacierEQ — Honolulu, HI

- Run an eval-harness discipline on every agent shipped: prompt and configuration changes validated by eval harnesses with replay loops before anything touches production — an eval loop defensible to a skeptic.
- Build guardrails into the orchestration layer: typed state transitions, approval gates on consequential actions, and fail-closed recovery — the defined 'what happens when it's unsure' path is halt-and-escalate, never guess.
- Monitor every run: per-task latency instrumentation, failure tracking, performance monitoring with A/B testing — no agent completion claim without independent readback.
- Track latency and token cost per agent run as first-class benchmarks — measured, logged, and priced into the eval discipline rather than discovered in production.
- Ship structured outputs, function calling with typed tool schemas, and retrieval-augmented generation over the estate across a **1,395-repository** platform with a **989-test** control plane — **62/62** verification checks green.
- Integrate **65 MCP connectors** plus APIs, webhooks, and data pipelines into environments I don't own — evidence boundaries and audit receipts built for client-cloud compliance regimes.

EDUCATION

University of Hawaii at Manoa 2016
Bachelor of Science

AWS Cloud Institute 2025
Cloud Computing program

Colorado State University
Undergraduate coursework

SELECTED PROOF

Job-App Helix — attributable campaign engine

- Composes assessments into campaign-level **GO / NO-GO** with immutable evidence models, stable finding codes, and fail-closed decisions.
- Recovery restricted to contingency evidence declared before reassessment — the engine never silently tunes values until they pass.

GlacierEQ estate — agent systems portfolio

- Durable patterns for claim tracking, evidence boundaries, audit receipts, and human-in-the-loop control across the full estate.
- Connects multiple technical surfaces — agents, APIs, databases — into one coherent, verifiable working system.

SKILLS

Languages: Python · TypeScript · SQL · Bash · JavaScript / Node.js

AI / Agents: LLM evaluation harnesses · replay loops · guardrails · hand-rolled multi-agent loops · function calling · structured outputs · RAG · MCP protocol · typed handoff contracts · prompt & context engineering

Data & Infra: Supabase · Postgres · Vercel · GitHub Actions · REST APIs · webhooks · containers

Practice: client-embedded delivery · demos & POCs · monitoring & observability · latency & token-cost benchmarks · verification gates · audit receipts · idempotent automation

Evidence boundary: portfolio repositories are independent work. Domain-inspired demonstrations are not claims of company affiliation, production deployment, or certification.