Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app

PROFILE

I build production-grade AI agents and the evaluation, observability, and guardrail frameworks that make them trustworthy: eval suites with golden sets and behavioral assertions, full trace logging, judge-in-the-loop rubrics, fail-closed gates, and human-in-the-loop where it matters. Security-first by construction — evidence boundaries, typed handoff contracts, approval gates on consequential actions, audit receipts on every run.

Three years as a founder running an evidence-grade agent control plane from Honolulu: hand-rolled multi-agent orchestration (planning, tool-use, retrieval, memory, guardrails, human handoffs) across a **1,395-repository** estate; **65** MCP connectors as the governed integration layer; per-task latency and token-cost instrumentation so agent cost is measured, not guessed; a **989-test** suite with **62/62** verification checks green.

CORE COMPETENCIES

Engineering: Python · TypeScript / Node.js · SQL · Bash · REST APIs · containers · GitHub Actions (CI/CD) · Supabase · Postgres · event-driven systems · async job queues

Applied AI / Agentic Systems: LLM-powered agents · multi-agent orchestration · tool use & structured outputs · MCP servers & connectors · RAG & reranking · evaluation frameworks · judge-in-the-loop · observability · guardrails · prompt engineering (Claude, GPT, Llama, Mistral) · model-agnostic benchmarking

Production & Security: end-to-end ownership (prototype → production) · fail-closed recovery · security-first design (permissions, data access, failure modes) · cost tracking & efficiency · CI regression gates · decision audit receipts · documentation & enablement

EXPERIENCE

Founder / Applied AI Systems Architect 2023 – Present

GlacierEQ — Honolulu, HI

- Design and build the eval, observability, and guardrail frameworks that decide whether an agent is actually doing the job: golden suites, behavioral assertions, judge rubrics, full trace logging, regression gates — measured against a baseline, every run, before anything touches production.
- Ship agents that take real actions under security-first design: evidence boundaries on every tool, typed handoff contracts between agents, human approval gates on consequential actions, fail-closed on ambiguous state — recovery restricted to contingency evidence declared before reassessment.
- Own multi-agent orchestration end to end: hand-rolled planning and tool-use loops with retrieval, memory, and guardrails; the durable patterns their Agentic Systems team is standardizing — reusable toolkits, shared best practices, and written enablement so every new agent ships on the same spine.
- Instrument agent cost like infrastructure: per-task latency and token-cost tracking across runs, model-agnostic benchmarking on every release — the usage-and-cost visibility the role asks for, built into the platform, not bolted on.

- Work the ambiguity the role describes: scope workflows with non-technical operators, turn loosely defined problems into agent-driven solutions, and state plainly what each agent can and cannot reliably do.
- Author **65** MCP connectors and integrations that expose tools and structured resources to any MCP-compatible system — ERPs, CRMs, data pipelines, webhooks, messaging — the shared, governed integration layer every downstream agent runs on.
- Run the control plane across a **1,395-repository** estate with a **989-test** suite, **62/62** verification checks green — no agent completion claim without independent provider readback.

EDUCATION

University of Hawaii at Manoa 2016
Bachelor of Science

AWS Cloud Institute 2025
Cloud Computing program

Colorado State University
Undergraduate coursework

SELECTED PROOF

Job-App Helix — attributable campaign engine

- Immutable evidence models, stable finding codes, and fail-closed **GO / NO-GO** decisions on an hourly production pipeline; no agent completion claim without independent provider readback.
- Reusable-pattern discipline: one-off field wins become shared infrastructure with versioned eval suites and written design docs — the 'reusable patterns, shared tools, best practices' motion.

GlacierEQ estate — production agent systems

- Durable patterns for claim tracking, evidence boundaries, audit receipts, and human-in-the-loop control across the full estate.
- Per-task latency instrumentation and decision audit logs across agent runs — the tracing practice that makes benchmarking honest.

SKILLS

Languages: Python · TypeScript / Node.js · SQL · Bash · Java

Applied AI: LLM agents · multi-agent orchestration · tool use & structured outputs · MCP integrations · RAG & reranking · evaluation frameworks · judge-in-the-loop · observability · guardrails · prompt engineering · model-agnostic benchmarking

Production & Security: end-to-end ownership · fail-closed recovery · security-first design · cost tracking & efficiency · CI regression gates · audit receipts · documentation & enablement

Evidence boundary: portfolio repositories are independent work. Domain-inspired demonstrations are not claims of company affiliation, production deployment, or certification.