Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app · linkedin.com/in/caseybartonai

PROFILE

Builder-first founding AI engineer. Three years running an evidence-grade agent control plane from Honolulu — I ship domain-specific agents into live operator workflows, then make them repeatable: evaluation infrastructure that measures whether agents actually replace human workflows (not vibes), model selection and routing across frontier + open models decided by measured economics, and AI-engineering playbooks every new engagement inherits on day one. 90% build is my natural ratio. Field fluency runs deeper than theory: I come from the physical economy — licensed home-inspector background, permits, site work, materials — so unglamorous industries are home turf.

My integration layer is **65** authored MCP connectors — governed tool/API/data-source connections with typed contracts and structured validation on every boundary — plus eval harnesses with written rubrics (LLM-as-a-judge), golden datasets, and behavioral assertions. Agents must be reliable, observable, and auditable from day one: the operating posture, not a compliance add-on. Honest on scope: no client-side enterprise deployments, no enterprise logos claimed; the building, the evals, and the playbooks are the transferable part.

CORE COMPETENCIES

Forward-deployed delivery: embedded delivery posture · customer discovery · sit-with-the-customer workflow instrumentation · ambiguous-scope scoping · rapid prototyping → production · legacy workflow replacement · handover discipline

Agentic AI: domain-specific agents (reason / plan / act) · tool use & function calling · structured outputs · MCP servers & connectors · RAG pipelines · multi-agent loops · guardrails · fail-closed recovery · human-in-the-loop gates

Evaluation, quality & model economics: eval infrastructure · golden datasets · behavioral assertions · LLM-as-a-judge with written rubrics · success-criteria framing (agents replacing workflows, measured) · frontier + open routing by cost/quality economics · trace intelligence

Production: Python (primary) · TypeScript · SQL · Bash · REST APIs · JSON Schema / Pydantic-style validation · GitHub Actions (CI/CD) · Postgres · per-task latency & token-cost instrumentation · retry logic

EXPERIENCE

Founder / Applied AI Systems Architect 2023 – Present
GlacierEQ — Honolulu, HI

- Ship working agents into live environments fast: the founding motion is domain-specific agent prototyped with the operator, measured against their real workflow, then hardened — 30-day-style delivery cadence across engagements.
- Build the eval infrastructure that measures whether agents replace human workflows — not vibes: golden datasets, behavioral assertions, LLM-as-a-judge with written rubrics, and trace intelligence feeding prompt/retrieval/configuration iteration.
- Decide fine-tune vs. prompt vs. RAG vs. agent vs. harness with numbers: model strategy is eval-driven selection and routing across frontier + open models by measured quality, latency, and cost — opinions

defended by the data, not the trend. (No weight fine-tuning claimed; the routing and selection economics are the proven part.)

- Sit with the operator, watch the workflow, instrument it, rebuild it: embedded scoping turns ambiguous field problems into well-framed agent deployments with clear success criteria.
- Make it repeatable: reusable AI-engineering playbooks, model stacks, eval harnesses, and integration patterns — every new engagement inherits the stack and the lessons, not a blank slate.
- Connect agents to the real world through **65** authored MCP connectors — governed integrations with typed contracts, evidence boundaries, and audit receipts on every tool boundary.
- Run an evidence-grade agent control plane across an **1398-repository** estate (live-verified 2026-10-20): observability, trace intelligence, recoverable runs — backed by a **989-test** suite, **62/62** verification checks green, and deterministic CI regression gates.

EDUCATION

University of Hawaii at Manoa **2016**
Bachelor of Science

AWS Cloud Institute 2025
Cloud Computing program

Colorado State University
Undergraduate coursework

SELECTED PROOF

Job-App Helix — evaluation-gated campaign engine

- Attributable campaign engine with immutable evidence models and fail-closed GO / NO-GO decisions — the eval gate decides what carries load; written rubrics behind every automated quality decision.

Governed integration layer — 65 authored MCP connectors

- Governed tool/API/data-source connections with typed contracts and structured validation on every boundary — the reusable integration layer behind every deployed agent engagement.

Failure posture — designed for production, not the demo: instrument the failure modes found in production — stuck loops, bad retrievals, tool misuse, context blowouts — and build the catchers: replay loops, fail-closed gates, recovery restricted to evidence declared before reassessment.

SKILLS

Languages: Python · TypeScript · SQL · Bash · JavaScript / Node.js

Agentic AI: domain-specific agents · tool/function calling · MCP · structured outputs · RAG · eval harnesses · LLM-as-a-judge · trace intelligence · guardrails · fail-closed recovery

Deployed engineering: discovery & scoping · workflow instrumentation · rapid prototyping · embedded delivery · field-to-product feedback · handover discipline · technical writing

Production: REST APIs · Postgres · CI/CD (GitHub Actions) · JSON Schema validation · observability · latency & token-cost instrumentation · retry logic

Evidence boundary: portfolio repositories are independent work. No enterprise customer logos claimed; client-side deployment depth is founder/operator (three years), honestly stated. No weight fine-tuning claimed — model strategy is eval-driven selection and routing. Home-inspector background = construction inspection — physical-economy fluency, not engineering depth. No Go / GCP / GraphQL / K8s / Terraform / DS depth claimed.