Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app

PROFILE

I own the production journey: lead AI workloads from proof-of-concept to production — deployed, stable, scalable — and drive time-to-production and time-to-value with customer engineering teams. The Token Factory motion (engineering × delivery × customer success) is my operating pattern: understand the customer's architecture and use cases, monitor and improve latency, throughput, cost, and reliability, resolve bottlenecks before they become incidents, and translate field evidence into structured product feedback.

Three years as a founder running an evidence-grade agent control plane from Honolulu: per-task latency and token-cost instrumentation on every run; failure-injection recovery and fail-closed gates as production discipline; **65** authored MCP connectors with typed contracts behind every tool boundary; structured feedback loops from field evidence to product iteration. Python primary, pragmatic polyglot, calm under incident pressure.

CORE COMPETENCIES

Customer delivery: PoC → production transition ownership · customer architecture assessment · latency / throughput / cost / reliability optimization · bottleneck identification · incident management posture · technical-customer communication · structured product feedback

AI production systems: inference-pipeline integration (client side) · AI workload evaluation · eval regression gates · behavioral assertions · golden datasets · trace intelligence · MCP servers & connectors · function calling · structured outputs · guardrails · fail-closed recovery · failure-injection testing

Production: Python (primary) · TypeScript / JavaScript · SQL · Bash · REST APIs · GitHub Actions (CI/CD) · Postgres · JSON Schema / Pydantic-style validation · retry logic · observability & run logging · per-task latency & token-cost instrumentation · documentation engineers actually read

EXPERIENCE

Founder / Applied AI Systems Architect 2023 – Present

GlacierEQ — Honolulu, HI

- Own the production journey: lead AI workloads from proof-of-concept to production — deployed, stable, scalable — and drive time-to-production and time-to-value with the teams that operate them.
- Ensure technical success in production: per-task latency and token-cost instrumentation on every run; trace intelligence naming the slowest calls, redundant tool use, and expensive segments; bottlenecks identified and resolved before they become incidents.
- Act as a trusted technical partner: work directly with technical counterparts, understand their architectures and use cases, provide guidance on best practices and optimization — and translate technical challenges into actionable solutions, clearly, under pressure.
- Manage risks and incidents like production work: fail-closed gates, behavioral assertions, recovery restricted to evidence declared before reassessment — calm communication and documented resolution paths when it matters.

- Drive continuous improvement: structured feedback from field evidence to product iteration — identify opportunities to optimize and expand usage, shape better solutions from real needs.
- Run an evidence-grade agent control plane across a **1,395-repository** estate: typed handoff contracts, guardrails, human approval gates on consequential actions — backed by a **989-test** suite, **62/62** verification checks green.

EDUCATION

University of Hawaii at Manoa 2016
Bachelor of Science

AWS Cloud Institute 2025
Cloud Computing program

Colorado State University
Undergraduate coursework

SELECTED PROOF

Job-App Helix — evaluation-gated campaign engine

- Attributable campaign engine with immutable evidence models and fail-closed **GO / NO-GO** decisions — the production-journey discipline: prove it works, then ship it.
- Field-to-product feedback loop in practice: trace and run evidence feeding configuration iteration — the structured feedback Token Factory runs on.

Inference-pipeline integration, honestly bounded

- Integration-side experience with inference pipelines (fal.ai + Whisper pipeline, estate-verified) — client-side integration, eval-gated, not GPU-infrastructure operations.
- Eval/trace engineering is the transferable production muscle: measure latency, cost, quality; gate on evidence; recover by design.

SKILLS

Languages: Python · TypeScript / JavaScript · SQL · Bash · Node.js

AI production: PoC → production ownership · inference-pipeline integration · eval regression gates · behavioral assertions · trace intelligence · latency/cost optimization · failure-injection recovery

Customer engineering: technical-customer delivery · incident communication · cross-team coordination · structured product feedback · handover discipline · technical writing

Integration: MCP servers & connectors · REST APIs · function calling · structured outputs · JSON Schema validation · guardrails · observability & run logging

Evidence boundary: portfolio repositories are independent work. No GPU-infrastructure, neocloud-operations, vLLM/TensorRT-production, or Kubernetes depth claimed — inference-framework knowledge is transferable (integration-side, eval/trace engineering). Customer-facing depth is founder/operator (three years), not enterprise accounts. No Nebius-internal claims. GitHub is the forge of record.