Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app · linkedin.com/in/caseybartonai

PROFILE

I own AI evaluation practice as a technical discipline, not a vibe: representative datasets, behavioral scenarios, scoring rubrics, regression suites, and human-review loops — with the eval harnesses, monitoring, and quality dashboards that make quality measurable without one-off manual checks. Three years as a founder running an evidence-grade agent control plane from Honolulu: I set quality baselines and fail-closed release gates for prompt, model, and agent changes, and I diagnose failures across prompts, context, orchestration, models, tools, and product code — then fix the right layer.

The tuning is eval-driven, never claimed weight work: **65** authored MCP connectors with typed contracts on every tool boundary, LLM-as-a-judge with written rubrics, golden datasets, and trace intelligence feeding prompt/retrieval/configuration iteration. Honest on scope: no weight fine-tuning claimed — eval-driven selection and routing is the proven part; founder/operator (three years, 2023–present), honestly stated.

CORE COMPETENCIES

AI evaluation practice: representative datasets · behavioral scenarios · scoring rubrics (LLM-as-a-judge) · regression suites · human review loops · quality baselines · release gates for prompt/model/agent changes · experimental design

Eval harnesses & observability: eval harness design & operation · monitoring & quality dashboards · trace intelligence · latency/cost/quality instrumentation · failure analysis across prompts, context, orchestration, models, tools

Agentic systems: agent loops & tools · structured outputs · multi-turn agents with tool use · user-intent preservation as behavioral assertions · MCP servers & connectors · guardrails · fail-closed recovery

Engineering: Python (primary) · TypeScript · SQL · prompt/instruction versioning with change records · GitHub Actions (CI/CD) · Postgres · REST APIs · JSON Schema validation

EXPERIENCE

Founder / Applied AI Systems Architect 2023 – Present
GlacierEQ — Honolulu, HI

- Own the evaluation practice end to end: define how the system should behave, write the quality definitions and scoring rubrics, and build the technical tooling — eval harnesses, regression suites, monitoring, quality dashboards — that make behavior measurable and the improvement loop continuous.
- Set quality baselines and release gates for prompt, model, and agent changes: every change carries a clear record of what changed, why it changed, and how behavior improved — fail-closed GO / NO-GO decisions when quality regresses.
- Diagnose failures across prompts, context, orchestration, models, tools, and product code: instrument the failure modes found in production — stuck loops, bad retrievals, tool misuse, context blowouts — and fix the right layer, with production failures and traces turned into durable regression cases.

• Combine automated checks with human judgment: **989-test** suite with behavioral assertions, **62/62** verification checks green, and human-review loops where judgment can't be automated — across an **1398-repository** estate (live-verified 2026-10-20).

• Connect agents to tools and data through **65** authored MCP connectors — governed connections with typed contracts, structured validation, and audit receipts on every boundary, so evaluation measures real tool behavior, not stubs.

• Tune with data, not opinion: model selection and routing across frontier + open models decided by measured quality, latency, and cost economics — eval-driven, never claimed weight fine-tuning.

• Translate ambiguous quality complaints into testable definitions: member-behavior ambiguity in, concrete system changes and regression coverage out — product sense expressed as written rubrics.

EDUCATION

University of Hawaii at Manoa **2016**
Bachelor of Science

AWS Cloud Institute 2025
Cloud Computing program

Colorado State University
Undergraduate coursework

SELECTED PROOF

Evaluation-gated quality — measured before it ships

• Evaluation harnesses with written rubrics (LLM-as-a-judge) gate every deployment: golden datasets, behavioral assertions, fail-closed gates — the release-gate discipline applied to every run.

Job-App Helix — evaluation-gated campaign engine

• Immutable evidence models with fail-closed GO / NO-GO decisions and per-run verification receipts — the continuous-improvement loop with the measurement built in.

Failure posture — diagnose the right layer

• Production failure modes instrumented (stuck loops, bad retrievals, tool misuse, context blowouts) with replay loops and recovery restricted to evidence declared before reassessment — durable regression cases, not one-off fixes.

SKILLS

Evaluation: golden datasets · behavioral scenarios · scoring rubrics · LLM-as-a-judge · regression suites · release gates · experimental design · human-review design

Agentic AI: agent loops · tool/function calling · MCP · structured outputs · RAG · guardrails · fail-closed recovery · trace intelligence

Engineering: Python · TypeScript · SQL · REST APIs · Postgres · CI/CD (GitHub Actions) · JSON Schema validation · observability · latency & token-cost instrumentation

Evidence boundary: portfolio repositories are independent work. Founder/operator depth (three years, 2023–present), honestly stated — no enterprise customer logos, no enterprise-account track record invented. No weight fine-tuning claimed — model strategy is eval-driven selection and routing. No Go / GCP / GraphQL / K8s / Terraform / data-science-statistics depth claimed. Eval-platform experience is self-built harnesses with comparable mechanics — Langfuse/Braintrust listed as bonus only; no depth on those platforms claimed.