Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app · linkedin.com/in/caseybartonai

PROFILE

Evaluation engineer for answer quality: I build the measurement systems that tell an AI product whether its answers are actually good — golden datasets, LLM-as-a-judge with written rubrics, behavioral assertions, and trace intelligence that turns production runs into the next eval iteration. Three years as a founder running an evidence-grade agent control plane where nothing ships until the evals say it should: fail-closed GO / NO-GO decisions, written rubrics behind every automated quality call. Answer quality is the product metric I optimize — measured, not vibed.

The systems are real: **65** authored MCP connectors feeding agents whose answer quality I measure daily; a **989-test** suite and **62/62** verification checks green across a **1398-repository** estate. Production systems shipped and operated — the 'shipping and operating production systems' bar, honestly stated at founder/operator scale.

CORE COMPETENCIES

Evaluation & answer quality: eval harnesses · LLM-as-a-judge with written rubrics · golden datasets · behavioral assertions · answer-quality metrics · regression gates · trace intelligence

Production ML systems: shipping and operating production systems · observability · latency & token-cost instrumentation · CI/CD (GitHub Actions) · incident-shaped thinking

Agentic AI: domain-specific agents · tool/function calling · MCP · structured outputs · RAG · guardrails · fail-closed recovery

Production: Python (primary) · TypeScript · SQL · Bash · REST APIs · JSON Schema / Pydantic-style validation · Postgres · retry logic

EXPERIENCE

Founder / Applied AI Systems Architect **2023 – Present**
GlacierEQ — Honolulu, HI

- Own answer quality end to end: define what 'good' means in written rubrics, build the judges and datasets that measure it, gate releases on the numbers — the eval decides what ships.
- Ship and operate production AI systems: agents in live operator workflows, monitored, iterated, and held to quality bars — three years of production ownership at founder/operator scale.
- Turn production into evals: trace intelligence over real runs feeds golden-dataset growth and rubric refinement — the measurement system improves with every deployment.
- Instrument the failure modes that degrade answers: stuck loops, bad retrievals, tool misuse, context blowouts — caught by behavioral assertions and fail-closed gates before users see them.
- Connect measurement to the real world through **65** authored MCP connectors — governed tool/API/data-source connections with typed contracts, structured validation, and audit receipts.

EDUCATION

University of Hawaii at Manoa **2016**
Bachelor of Science

AWS Cloud Institute 2025
Cloud Computing program

Colorado State University
Undergraduate coursework

SELECTED PROOF

Evaluation-gated quality — the eval decides what ships

• Fail-closed GO / NO-GO campaign engine: immutable evidence models, per-run verification receipts — the discipline of declared success criteria applied to every production engagement.

Answer-quality machinery — rubrics, judges, datasets

• LLM-as-a-judge with written rubrics, golden datasets, behavioral assertions: directly shapes how answer quality is measured and improved — the posting's core sentence, as built systems.

SKILLS

Languages: Python · TypeScript · SQL · Bash · JavaScript / Node.js

Eval & quality: eval harnesses · LLM-as-a-judge · written rubrics · golden datasets · behavioral assertions · regression gates · trace intelligence · failure analysis

Production: REST APIs · MCP · Postgres · CI/CD (GitHub Actions) · JSON Schema validation · observability · latency & token-cost instrumentation

Evidence boundary: portfolio repositories are independent work. Three years stated as three years — against the posting's 4+ year bar, honestly stated; the eval/production-ownership depth is the case. No weight fine-tuning claimed. No Perplexity-internal claims; transferable eval/quality engineering only. No Go / GCP / GraphQL / K8s / Terraform / DS-statistics depth claimed.