Casey Barton

Honolulu, HI · (808) 936-5654 · glacier.equilibrium@gmail.com · github.com/GlacierEQ · casey-barton-glaciereq.vercel.app · linkedin.com/in/caseybartonai

PROFILE

Tech-lead posture from a founder's seat: I own AI-native domains outright — agent architecture, platform abstractions, evaluation and guardrail infrastructure — setting technical direction, writing the load-bearing design docs, and staying in the code for the hardest parts. Three years shipping production agentic systems from Honolulu: agents with retrieval, orchestration, policy-based routing, tool invocation, evaluation harnesses, and lifecycle observability — production-ready standards held in review: testability, rollback safety, cost ceilings, failure modes, human-in-the-loop boundaries.

My platform layer is **65** authored MCP connectors — governed tool/API/data-source connections with typed contracts, structured validation, and audit receipts — plus eval infrastructure (golden datasets, LLM-as-a-judge with written rubrics, regression gates) and reusable patterns every engagement inherits. Honest on scope: no weight fine-tuning claimed — eval-driven selection and routing is the proven part; founder/operator (three years, 2023–present), honestly stated. No regulated-environment (HIPAA) production claimed — only the engineering posture transfers: data minimization, access controls, audit logging, and staged release discipline as architecture, not afterthoughts.

CORE COMPETENCIES

Technical leadership: domain ownership · tech-lead decomposition & sequencing · design docs · load-bearing architectural calls · ambiguity removal · cross-team influence without positional authority · mentorship by review

Agent architecture: production AI agents · retrieval · orchestration · policy-based routing · tool invocation · evaluation harnesses · guardrail infrastructure · lifecycle observability · human-in-the-loop boundaries

AI platform foundations: platform abstractions & shared libraries · 65 authored MCP connectors · typed contracts · reusable engineering patterns · privacy/security/compliance as architecture

Cloud-native & full-stack: TypeScript/Node.js + Python · SQL (relational fluency) · REST APIs · CI/CD (GitHub Actions) · observability · on-call-ready operations · latency & token-cost instrumentation

EXPERIENCE

Founder / Applied AI Systems Architect 2023 – Present
GlacierEQ — Honolulu, HI

- Own the technical direction for an AI-native domain end to end: agent architecture, platform abstractions, and evaluation/guardrail infrastructure across an **1398-repository** estate (live-verified 2026-10-20) — decomposing ambiguous problems, sequencing delivery, clearing blockers, keeping the work pointed at the outcome rather than the ticket.
- Ship production AI agents incorporating retrieval, orchestration, policy-based routing, tool invocation, evaluation harnesses, and lifecycle observability — agents that carry real operational load, measured after the fact, not just at demo time.
- Set the production-ready bar and hold it: testability, rollback safety, cost ceilings, failure modes, human-in-the-loop boundaries — **989-test** suite, **62/62** verification checks green, deterministic CI regression gates, fail-closed recovery on every deployed agent.

• Build the abstraction layers teams build on: **65** governed MCP connectors plus shared patterns, typed contracts, and structured validation — the reusable platform surface behind every deployed engagement.

• Own the evaluation strategy per domain: define the metrics, test harnesses, and evaluation plans that measure agent accuracy, latency, safety, and cost-effectiveness — kill approaches that don't work, early and visibly.

• Design for failure and advise as a technologist: instrument production failure modes — stuck loops, bad retrievals, tool misuse, context blowouts — with the catchers (replay loops, fail-closed gates, behavioral assertions), and run design sessions alongside operators and stakeholders so field friction becomes engineering requirements that move the roadmap.

EDUCATION

University of Hawaii at Manoa **2016**
Bachelor of Science

AWS Cloud Institute 2025
Cloud Computing program

Colorado State University
Undergraduate coursework

SELECTED PROOF

Governed integration layer — 65 authored MCP connectors

• Typed contracts and structured validation on every tool boundary — the platform-abstraction discipline, including enterprise-style identity/access patterns: approval gates, access controls, audit receipts.

Evaluation-gated quality — the measure-and-improve loop

• Golden datasets, behavioral assertions, LLM-as-a-judge with written rubrics, trace intelligence — the evaluation story behind every shipped agent; fail-closed GO / NO-GO decisions keep regression from carrying load.

SKILLS

Languages: Python · TypeScript / Node.js · SQL · Bash · JavaScript

Agentic AI: agent workflows · retrieval · orchestration · policy-based routing · tool/function calling · MCP · structured outputs · eval harnesses · guardrails · lifecycle observability

Platform & production: platform abstractions · shared libraries · CI/CD (GitHub Actions) · Postgres · observability · retry logic · latency & token-cost instrumentation · design docs

Evidence boundary: portfolio repositories are independent work. Founder/operator depth (three years, 2023–present), honestly stated — no enterprise customer logos, no enterprise-account track record invented. No weight fine-tuning claimed — model strategy is eval-driven selection and routing. No Go / GCP / GraphQL / K8s / Terraform / data-science-statistics depth claimed. Founder/operator depth is three years, 2023–present — stated against the JD's 8+ years (see REVIEW_REQUIRED.txt). No regulated/HIPAA production claimed; compliance-as-architecture is engineering practice, not domain experience.